A3: Computer Networking (I)

Computer Systems 2024-25
Department of Computer Science
University of Copenhagen

David Marchant

Due: Sunday, 17th of November, 16:00
Version 1 (October 25, 2024)

This is the fourth assignment in the course on Computer Systems 2022 at DIKU and the first on the topic of Computer Networks. We encourage pair programming, so please form groups of 2 or 3 students. Groups cannot be larger than 3, and we strongly recommend that you do not work alone.

For this assignment you will receive a mark out of 4 points; You must attain at least half of the possible points to be admitted to the exam. For details, see the *Course description* in the course page. This assignment belongs to the CN category together with A4. Resubmission is not possible.

It is important to note that only solving the theoretical part will result in 0 points. Thus, the theoretical part can improve you score, but you *have* to give a shot at the programming task.

The web is more a social creation than a technical one. I designed it for a social effect — to help people work together — and not as a technical toy. The ultimate goal of the Web is to support and improve our weblike existence in the world. We clump into families, associations, and companies. We develop trust across the miles and distrust around the corner.

— Tim Berners-Lee, Weaving the Web (1999)

Overview

This assignment has two parts, namely a theoretical part (Section 1) and a programming part (Section 2). The theoretical part deals with questions that have been covered by the lectures. The programming part requires you to complete a file transfer service using socket programming in C. In this assignment, the programming effort relies solely on building the client portion of the architecture. More details will follow in the programming part (Section 2).

1 Theoretical Part (15%)

You should also hand-in answers to the following questions. These questions are each quite open ended, and unless otherwise noted should take only a paragraph or two to answer.

1.1 Reserving Connections

At the start of the networking section we introduced circuit-switching which reserves a connection across the network. We explained that this was inefficent and instead the modern internet uses a connectionless packet-based system instead. Later on we explained how most communication across the network is done via TCP, which is connection oriented. How is it true that our modern internet is both connectionless and connection dependent?

1.2 Hashing

Rainbow tables only work if we hash inputs using known hash functions. Rather than bothering with salts, an alternate solution could be to use a hash function uniquely developed for each application. Why is this not advised? You should explain what features a hash function is expected to have, and how this approach might compromise those features, as well as any other additional relevent points.

1.3 Password salting

The protocol described at the end of this document uses a salt on both the client and server side of the communication. One limitation of the current system is that the client generates the salt for its own use, and therefore if the same user logs onto the system from a new client they will be told they cannot log on, even if they use the same username and password.

Design a protocol that would support the client-side salt being shared between multiple clients. It is up to you to determine who generates the salt, and to define the format for any new message types that will be sent over the network, along with any appropriate responses if so required. You should also provide a brief description of how clients and servers may behave as the systems progresses.

You can use the existing protocol as a basis, though some parts will need to be added/modified. A complete answer for this may stretch over several pages depending on how much you add or duplicate from the existing protocol, though a good answer will make the minimum necessary additions/modifications. It may be easier to answer this question after you have completed the programming section below and are more familiar with the existing protocol.

Note that you do not need to implement any code for this question, and your handed in code for the programming part of this assignment does not need to implement this functionality.

2   Programming Part (85 %)

For the programming part of this assignment, you will implement a client in a networked file-sharing service. Many of the programming tasks set in the exercises relate to the sub-tasks within this assignment, so if you are stuck at any point consider revisiting them.

2.1   Design and overview

The file-sharing service consists of server and client processes. For this assignment we can assume only one of each, with a single server providing files, and a single peer retrieving them.

This design uses a custom protocol outlined below. It has *some* security features built in, which you should use appropriately. By this it is meant that you should be able to explain how your implementation keeps to these specifications.

In this protocol, each client should start by registering with the server. This is done by simply telling the server the client is here, which is not a secure methodology in the real world but will suffice for this assignment. If a client attempts any further communication with the server without registering the user, then the server will refuse to provide any requested files. Throughout this protocol we will make repeated use of hashes. These are used for two broad purposes, to obfuscate passwords and login details, and to act as a checksum on our transported data.

The server is limited in how many bytes it can send at once. If the client requests a file larger than this limit then it will therefore be sent over several messages which the client will have to reassemble. There is no guarantee that these blocks are sent in any particular order.

For this assignment you are going to implement a client that registers a new user, and retrieves files and stores them locally.

2.2   API and Functionality

2.2.1   Key Implementation Tasks

The client you will be implementing needs to perform the six basic tasks listed below. Note that each of these tasks do not need to be completed in the order listed. Each should be completed according to the protocol described in Section 3. The client should:

2.1. Create a user signature. This is obtained by salting the user-remembered password and hashing it. This will be used by both the registration and file request messages so the result should either be saved somehow, or be computed in a deterministic manner. Initially you may wish to use a hard-coded hash.

2.2. Register a new user. Both the username and the signature should be sent to the server to register a new user, with the client reporting back the response.

2.3. Retrieve a small file. The user should request a file small enough to fit into a single message. For this the example file tiny.txt has been provided, though you are free to test on others. This should be retrieved by the client and written into its own file storage. The client should validate the recieved file to check that it has been recieved correctly.

2.4. Retrieve a large file. The user should request a file large enough to require several messages to completely send. For this the example file hamlet.txt has been provided, though you are free to test on others. This should be retrieved by the client and written into its own file storage. The client should validate the recieved file to check that it has been recieved correctly. Note that the chunks may be recieved out of order, and it is the job of the client to reassemble them into the correct ordering.

2.5. Randomly generate a salt for each user and save it between sessions (e.g. saved when a client is closed and reloaded next time it started) so that the same user uses the same salt. No example code has been handed out for this as it is up to determine a robust way to do this, and what security precautions you need to make. Your solution should support multiple users, each with a unique salt so you will need to design a protocol for saving salts in such a way that a users salt can be retrieved from a collection of user salts. For this assignment you can save the salt locally and ignore any issues from logging on with the same username and password from a seperate device.

2.6. Implement some user interaction. The user of this system should be capable of taking user inputs as to what files should be retrieved. This should be via typing into stdin, though it *may* also be worth implementing a version that can use arguments as this may assist with testing. The system should be robust to different output and always present the user with useful feedback.

2.2.2 Test environment

One of the difficulties with developing a system that relies on network communication is that besides a network, you need to have both the client and server running, and potentially debug both simultaneously.

To assist in developing and testing your code, we have provided implementations of the client and server, but written in Python. You can use these to get started with a setup that runs fully contained on your own machine. By running a local server and a local client it is possible to debug and monitor all pieces of the communication, which is often not possible with an existing system.

You can of course also peek into or alter the Python client and get inspiration for implementing you own client, but beware that what is a sensible design choice in Python *may* not be a good choice for C. You are also reminded that although you can alter the Python as much as you want to provide additional debugging or the like, they are currently correct implementations of the

defined protocols. Be careful in making changes that do not alter the impleme-
nation details, as this may lead your C implementation astray.

To start an appropriate test environment you can run the server using the
command below, run from the *python/server* directory:

```
python3 server.py ../../config.yaml
```

Assuming you have not altered `config.yaml`, this will start a server at
0.0.0.0:12345. This will be capable of serving any files within the *python/server*
directory. You can then start an ideal client by running another command
prompt from the *python/client* directory:

```
python3 client.py ../../config.yaml
```

This will start a client capable of user interaction. It will ask for a username
and password, register that user and then ask for input as to what files to re-
quest. Alternatively, you can start the client as a linear script that will not ask
for user input. It will generate a hard coded user and retrieve two files. This
has been included to demonstrate how you *might* wish to implement some au-
tomatic testing. It can be started with:

```
python3 client.py ../../config.yaml test
```

Both client types with host at 0.0.0.0:23456 by default and will write any
retrieved files to the *python/client* directory.

The C client can be started by calling the following from within the *src* di-
rectory:

```
make
./networking ../config.yaml
```

By default this will compile your client, though as not yet implemented it
will not do anything.

2.3   Run-down of handed-out code and what is missing

- *src/networking.c* contains the client, and is where it is expected all of your
  coding will take place. If you alter any other files, make sure to highlight
  this in your report. This code contains three guide functions, `get_signature`,
  `register_user` and `get_file` to illustrate where you might begin, but
  feel free to make whatever alterations to the structure you wish. No
  guiding functions have been added for the user interaction or saving salt
  tasks, it is up to you and your reading of the code to determine where
  best to implement them. Note that lines 193-212 are included as a way
  of running a client as a linear script, and should be removed as part of
  implementing user interaction. Your completed system should write any
  retrieved files to the *src* directory.

- *src/networking.h* contains a number of structs you may find useful in build-
  ing your client. You are not required to use them as is, or at all if you
  would rather solve problem some other way.

- *src/sha256.c* and *src/sha256.h* contain an open-source stand-alone implementation of the SHA256 algorithm, as used throughout the network. You should not need to alter this file in any way.

- *src/common.c* and *src/common.h* contains some helper functions to assist in parsing the configuration file. You should not need to alter this file in any way.

- *src/compsys_helpers.c* and *src/compsys_helpers.h* contains definitions for robust read and write operations. You should not need to alter this file in any way.

- *src/endian.h* contains definitions for converting between big and small endian numbers. This is only useful on a mac, and should not be necessary in the majority of solutions so do not be concerned if you do not use any of these functions. You should not need to alter this file in any way.

- *python/server/server.py* contains the server, written in Python. This can be run on Python 3.7 or newer, and requires a config file written in the YAML format. You should not need to alter this file, though may find adding more debug print statements helpful.

- *python/client/client.py* contains the client, written in Python. This can be run on Python 3.7 or newer, and requires a config file written in the YAML format. You should not need to alter this file, though may find adding more debug print statements helpful.

- *python/shared.py* contains some functions shared between the server and client. You should not need to alter this file, though may find adding more debug print statements helpful.

- *config.yaml* contains the definitions for where the server and client will be hosted.

- *python/server/tiny.txt* and *python/server/hamlet.txt* are example data file which if you can transfer to the client correctly are enough to say you have completed this part of the assignment. Testing using addition files would be advantagous in a report.

- `README` contains the various commands you will need to get this project up and running.

2.4 Testing

For this assignment, it is *not* required of you to write formal, automated tests, but you *should* test your implementation to such a degree that you can justifiably convince yourself (and thus the reader of your report) that each API functionality implemented works, and are able to document those which do not.

Simply running the program, emulating regular user behaviour and making sure to verify the result file should suffice, but remember to note your results. Just testing using the three function calls (e.g. register a correct user,

retrieve a small file, retrieve a large file) is not enough to say the system works. Other tests such as retrieving without first registering, or requesting a file that does not exist should also be run at a minimum.

One final resource that has been provided to assist you is that an instance of server.py has been remotely hosted, and which you can connect to. This is functionally identical to those contained in the handout, with the obvious difference that any communications with it will be propperly over a network.

The server is designed to be robust, so should be resilient to any malformed messages you send, but as it is only a single small resource be mindful of swamping it with requests and only use them once you are confident in your system. You can connect to the tracker at the following address and port:

```
Test Server: 130.225.104.138:5555
```

Note that depending on whatever firewalls or other network configuration options you have, you may not be able to reach the server from home, but that you *should* be able to from within the university.

2.5 Recommended implementation progress

As mentioned in the previous section, `src/networking.c` contains two unimplemented functions, the implementation of which should yield a functioning client. Do note that you are not limited to completing these functions and you are free to make whatever changes, additions or removals you wish at any point in the provided C code, as long as you successfully implement the described protocols. It is expected you will stick to modifying network.c, and should specifically highlight in your report if you deviate from this.

In this section, we give a short recap of your implementation todos; this can serve as a checklist for your project, and *we recommend* that you work on them in the order presented here.

2.1. Assemble the user signature from a password and salt.

2.2. Register a user with the server and print any feedback to the registration request.

2.3. Randomly generate a salt, and save it locally between client sessions.

2.4. Request a small file from the server and either create a local copy of the file, or print any feedback if an error is encountered.

2.5. Request a large file from the server and either create a local copy of the file, or print any feedback if an error is encountered.

2.6. Implement user interaction with your client. It should be possible for a user to request different files without restarting the client.

2.7. Manually test your implementation, documenting bugs found and how you fix them (if you are able to).

2.8. Meanwhile, do not neglect the theoretical questions. They may be relevant to your understanding of the implementation task.

2.6   Report and approximate weight

The following approximate weight sums to 85% and includes the implementation when relevant.

Please include the following points in your report:

- Discuss the provided protocol and how you were able to use it. For instance, what security assurances were you able to make using it? What gaps exist in this security setup and what changes, improvements or expansions should be made to the protocol to help address them? You may also wish to discuss how you save and load user salts, and how you do so according to whatever protocol you have designed. (Approx. weight: 20%)

- Document technical implementation you made for the client - cover in short each of the three tasks (register, small file, large file), as well as any additional changes you made. Each change made should be breifly justified. Make sure especially to discuss how you saved your salts between sessions, justifying your implementation. You should also describe what input/data your implementation is capable or processing and what it is not. (Approx. weight: 30%)

- Discuss how your design was tested. What data did you use on what machines? Did you automate your testing in any way? Remember that you are not expected to have built a perfect solution that can manage any and all input, but you are always meant to be able to recongise what will break your solution. Testing on the provided two files (tiny.txt and hamlet.txt) is not sufficient, you should also test using data that does not work, or with requests that are incorrect. (Approx. weight: 25%)

- Discuss any shortcomings of your implementation, and how these might be fixed. It is not necessarily expected of you to build a fully functional client, but it *is* expected of you to reflect on the project. Even if you implement the protocol exactly as described through this document there will still be problems to identify in terms of usability, security or functionality. (Approx. weight: 10%)

As always, remember that it is also important to document half-finished work. Remember to provide your solutions to the theoretical questions in the report pdf.

3   Protocol description

This section defines the implementation details of the components used in the A3 client-server file-sharing network.

3.1   File structures of the server and client

Both the server and client will be able to serve/write files relative to their own locations. In other words, any files in the same directory as the client or server should be servable.

3.2   Security considerations

In order to protect user-remembered passwords, they are always salted and then hashed. This is refered to as a signature within this protocol. The signature is treated effectively as a password by the system at large. The user-remembered password should NEVER be transmitted across the network.

The server will register a user by username and signature. The server will save this signature by salting it again, and hashing. This will then be saved along with the salt used by the server.

3.3   Format of the client requests

All messages from the client to the server must contain the same header as follows:

```
16 bytes - Username, as UTF-8 encoded bytes
32 bytes - Signature, a hash of the salted user password,
           as UTF-8 encoded bytes
 4 bytes - Length of request data, excluding this header,
           unsigned integer in network byte-order
```

If the message is intended as a registration of a new user then no additional data is provided, and so the length of the request data will be zero.

If the message is intended as a request for a file, then an additional field of the length given in the header will also be present. This is a path to the requested file to be retrieved.

Regardless of the message type, a response will ALWAYS be expected from the server.

3.4   Format of the server responses

All replies from the server to the client must contain the same header as follows:

```
 4 bytes - Length of response data, excluding this header,
           unsigned integer in network byte-order
 4 bytes - Status Code of the response,
           unsigned integer in network byte-order
 4 bytes - Block number, a zero-based count of which block in
           a potential series of replies this is,
           unsigned integer in network byte-order
 4 bytes - Block count, the total number of blocks to be sent,
           unsigned integer in network byte-order
32 bytes - Block hash, a hash of the response data in this
            message only, as UTF-8 encoded bytes
32 bytes - Total hash, a hash of the total data to be sent
            across all blocks, as UTF-8 encoded bytes
```

In addition to the header each response will include an additional payload of the length given in the header. Note that in the case of a small enough response to only require a single reply, then the block hash and total hash will be identical. If the response is to a request for data file, and it could be retrieved with no errors then the payload will be either all or part of said file, depending on the size of the file. In all other cases the response will be a feedback message explaining any errors or results of other querries.

The following error codes may be provided by the server:

```
1 - OK (i.e. no problems encountered)
2 - User already exists (i.e. could not register a user as
    they are already registered)
3 - User missing (i.e. could not service the request as the
    user has not yet registered)
4 - Invalid Login (i.e. the provided signature does not match
    the registered user)
5 - Bad Request (i.e. the request is coherent but cannot be
    served as the file doesn't exist)
6 - Other (i.e. any error not covered by the other status
    codes)
7 - Malformed (i.e. the request is malformed and could not be
    processed)
```

3.5   Requests/Responses and persistent connections

Note that there are no persistent connections between client requests. Each time a client registers or requests a file, the server will respond with one or more reply messages, but it will then close the connection. To request multiple files, a client will need to open a separate connection each time.

Submission

You should hand in a ZIP archive containing a `src` directory containing all *relevant* files (no ZIP bomb, no compiled objects, no auxiliary editor files, etc.).

Any Python code should also be included in your submission as it may form part of how you tested your C code (hint). As this course is not a Python course, you will not be marked according to the quality of your Python code.

To `make` this a breeze, we have configured the `src/Makefile` such that you can simply execute the following shell command within the `src` directory to get a `zip`:

```
$ make zip
```

Note that depending on any test files, or how you implemented your solutions you may wish to include other files not zipped by the above command. You are allowed to do so, but are encouraged to be sensible in what you include (E.g. please no 10GB test files).

Alongside a `src.zip` submit a `report.pdf`. For submission, sign up for a group at:

`https://absalon.ku.dk/courses/70194/groups#tab-23701`
Hand-in as a group as in other courses.