

Computer Systems: Network Programming (Sockets)

David Marchant

Based on slides by Randal E. Bryant and David R. O'Halloran, with alterations by Vivek Shah

But first, Unix I/O

- A Linux *file* is a sequence of m bytes:
 - $B_0, B_1, \dots, B_k, \dots, B_{m-1}$
- Cool fact: All I/O devices are represented as files:
 - `/dev/tty` (the current terminal)
 - `/dev/sda2` (a disk partition)
 - `/dev/tty2` (some other terminal)
- Even the kernel is represented as a file:
 - `/boot/vmlinuz-3.13.0-55-generic` (kernel image)
 - `/proc` (process information)
 - `/sys` (kernel data structures)


File Types

- Each file has a *type* indicating its role in the system
 - *Regular file*: Contains arbitrary data
 - *Directory*: Index for a related group of files
 - *Socket*: For communicating with a process on another machine

- Other file types beyond our core scope
 - *Named pipes (FIFOs)*
 - *Symbolic links*
 - *Character and block devices*

Unix I/O Overview

- Elegant mapping of files to devices allows kernel to export simple interface called *Unix I/O*:
 - Opening and closing files
 - `open()` and `close()`
 - Reading and writing a file
 - `read()` and `write()`
 - Changing the **current file position** (seek)
 - indicates next offset into file to read or write
 - `lseek()`
 - Not all files support seeking (e.g. pipes, sockets)



Current file position = k

Opening Files

- Opening a file informs the kernel that you are getting ready to access that file

```
int fd;    /* file descriptor */  
  
if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {  
    perror("open");  
    exit(1);  
}
```

- Returns a small identifying integer *file descriptor*
 - `fd == -1` indicates that an error occurred
- Each process created by a Linux shell begins life with three open files associated with a terminal:
 - 0: standard input (stdin)
 - 1: standard output (stdout)
 - 2: standard error (stderr)

Closing Files

- Closing a file informs the kernel that you are finished accessing that file

```
int fd;      /* file descriptor */
int retval;  /* return value */

if ((retval = close(fd)) < 0) {
    perror("close");
    exit(1);
}
```

- Closing an already closed file is a recipe for disaster in threaded programs, because the file descriptor number may have been re-used
- Always check return codes, even for seemingly benign functions such as `close()`

Reading Files

- Reading a file copies bytes from the current file position to memory, and then updates file position

```
char buf[512];
int fd;          /* file descriptor */
int nbytes;      /* number of bytes read */

/* Open file fd ... */
/* Then read at least 1 byte and
   up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

- Returns number of bytes read from file fd into buf
 - Return type `ssize_t` is signed integer
 - `nbytes < 0` indicates that an error occurred
 - **Short counts** (`nbytes < sizeof(buf)`) are possible and are not errors!

Writing Files

- Writing a file copies bytes from memory to the current file position, and then updates current file position

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open the file fd ... */
/* Then write up to 512 bytes from buf to file fd */
if ((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

- Returns number of bytes written from buf to file fd
 - **nbytes** < 0 indicates that an error occurred
 - As with reads, short counts are possible and are not errors!

Simple Unix I/O example

- Copying stdin to stdout, one byte at a time

```
int main(void)
{
    char c;

    while(read(STDIN_FILENO, &c, 1) != 0)
        write(STDOUT_FILENO, &c, 1);
    exit(0);
}
```

A Programmer's View of the Internet

1. Hosts are mapped to a set of 32-bit *IP addresses*
 - 128.2.203.179
2. The set of IP addresses is mapped to a set of identifiers called Internet *domain names*
 - 128.2.217.3 is mapped to `www.cs.cmu.edu`
3. A process on one Internet host can communicate with a process on another Internet host over a *connection*
4. This is **very similar** to the pipes we briefly introduced in the concurrency lectures.

Global IP Internet (upper case)

- Most famous example of an internet
- Based on the TCP/IP protocol family
 - IP (Internet Protocol) :
 - Provides *basic naming scheme* and unreliable *delivery capability* of packets (datagrams) from *host-to-host*
 - UDP (Unreliable Datagram Protocol)
 - Uses IP to provide *unreliable* datagram delivery from *process-to-process*
 - TCP (Transmission Control Protocol)
 - Uses IP to provide *reliable* byte streams from *process-to-process* over *connections*
- Accessed via a mix of Unix file I/O and functions from the *sockets interface*

IP Addresses

- 32-bit IP addresses are stored in an *address struct*
 - IP addresses are always stored in memory in *network byte order* (big-endian byte order)
 - True in general for any integer transferred in a packet header from one machine to another.
 - E.g., the port number used to identify an Internet connection.

```
/* Internet address structure */  
struct in_addr {  
    uint32_t    s_addr; /* network byte order (big-endian) */  
};
```

Dotted Decimal Notation

- By convention, each byte in a 32-bit IP address is represented by its decimal value and separated by a period
 - IP address: `0x8002C2F2` = `128.2.194.242`
- Use `inet_ntop`, `inet_pton` functions for converting between dotted decimal notation and IP addresses
 - Use `htonl`, `htons`, `ntohl` and `ntohs` functions for network byte order conversions
- Use `getaddrinfo` and `getnameinfo` functions (described later) to convert between IP addresses and dotted decimal format.

Internet Connections

- Clients and servers communicate by sending streams of bytes over *connections*. Each connection is:
 - *Point-to-point*: connects a pair of processes.
 - *Full-duplex*: data can flow in both directions at the same time,
 - *Reliable*: stream of bytes sent by the source is eventually received by the destination in the same order it was sent.
- A *socket* is an endpoint of a connection
 - *Socket address* is an **IPaddress:port** pair
- A *port* is a 16-bit integer that identifies a process:
 - **Ephemeral port**: Assigned automatically by client kernel when client makes a connection request.
 - **Well-known port**: Associated with some *service* provided by a server (e.g., port 80 is associated with Web servers)

Sockets

- What is a socket?
 - To the kernel, a socket is an endpoint of communication
 - To an application, a socket is a file descriptor that lets the application read/write from/to the network
 - **Remember:** All Unix I/O devices, including networks, are modeled as files
- Clients and servers communicate with each other by reading from and writing to socket descriptors




- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors

Socket Address Structures

- Generic socket address:
 - For address arguments to **connect**, **bind**, and **accept**
 - Necessary only because C did not have generic (**void ***) pointers when the sockets interface was designed
 - For casting convenience, we adopt the Stevens convention:
typedef struct sockaddr SA;

```
struct sockaddr {  
    uint16_t  sa_family;    /* Protocol family */  
    char      sa_data[14]; /* Address data. */  
};
```

sa_family



Family Specific

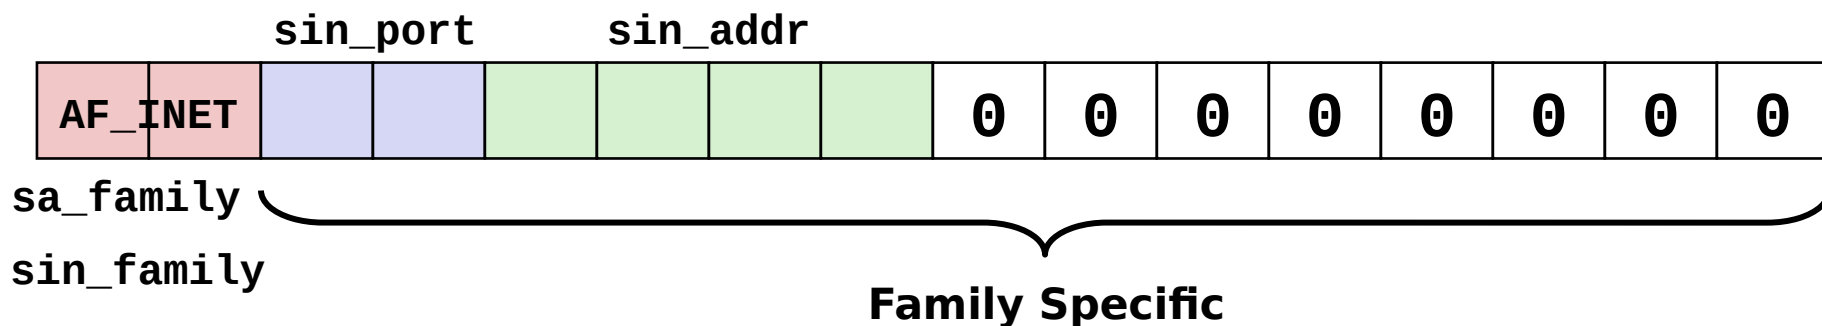
Socket Addresses Domains

- There are many different address domains, each with a corresponding struct.
- See link for complete list:
<https://man7.org/linux/man-pages/man2/socket.2.html>
- In this course we will stick with IPv4 (explained in the following slides)

Socket Address Structures

- Internet (IPv4) specific socket address:
 - Must cast (`struct sockaddr_in *`) to (`struct sockaddr *`) for functions that take socket address arguments.

```
struct sockaddr_in {
    uint16_t    sin_family; /* Protocol family (always AF_INET) */
    uint16_t    sin_port;   /* Port num in network byte order */
    struct in_addr sin_addr; /* IP addr in network byte order */
    unsigned char sin_zero[8]; /* Pad to sizeof(struct sockaddr) */
};
```



Setting up an address example

```
struct sockaddr_in serv_addr;  
  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_port = htons(12345);  
  
if (inet_pton(AF_INET, "123.123.123.123", &serv_addr.sin_addr) <= 0) {  
    printf("Invalid address\n");  
    return -1;  
}
```

- Note that we're already validating results. This is important in any application, but in networking its even more necessary than that
- Also note byte order on the port, and notation for the host

Host and Service Conversion: `getaddrinfo`

- **`getaddrinfo`** is the modern way to convert string representations of hostnames, host addresses, ports, and service names to socket address structures.
 - Replaces obsolete `gethostbyname` and `getservbyname` funcs.
- Advantages:
 - Reentrant (can be safely used by threaded programs).
 - Allows us to write portable protocol-independent code
 - Works with both IPv4 and IPv6
- Disadvantages
 - Somewhat complex
 - Fortunately, a small number of usage patterns suffice in most cases.

Host and Service Conversion: getaddrinfo

```
int getaddrinfo(const char *host,           /* Hostname or address */
               const char *service,        /* Port or service name */
               const struct addrinfo *hints, /* Input parameters */
               struct addrinfo **result);   /* Output linked list */

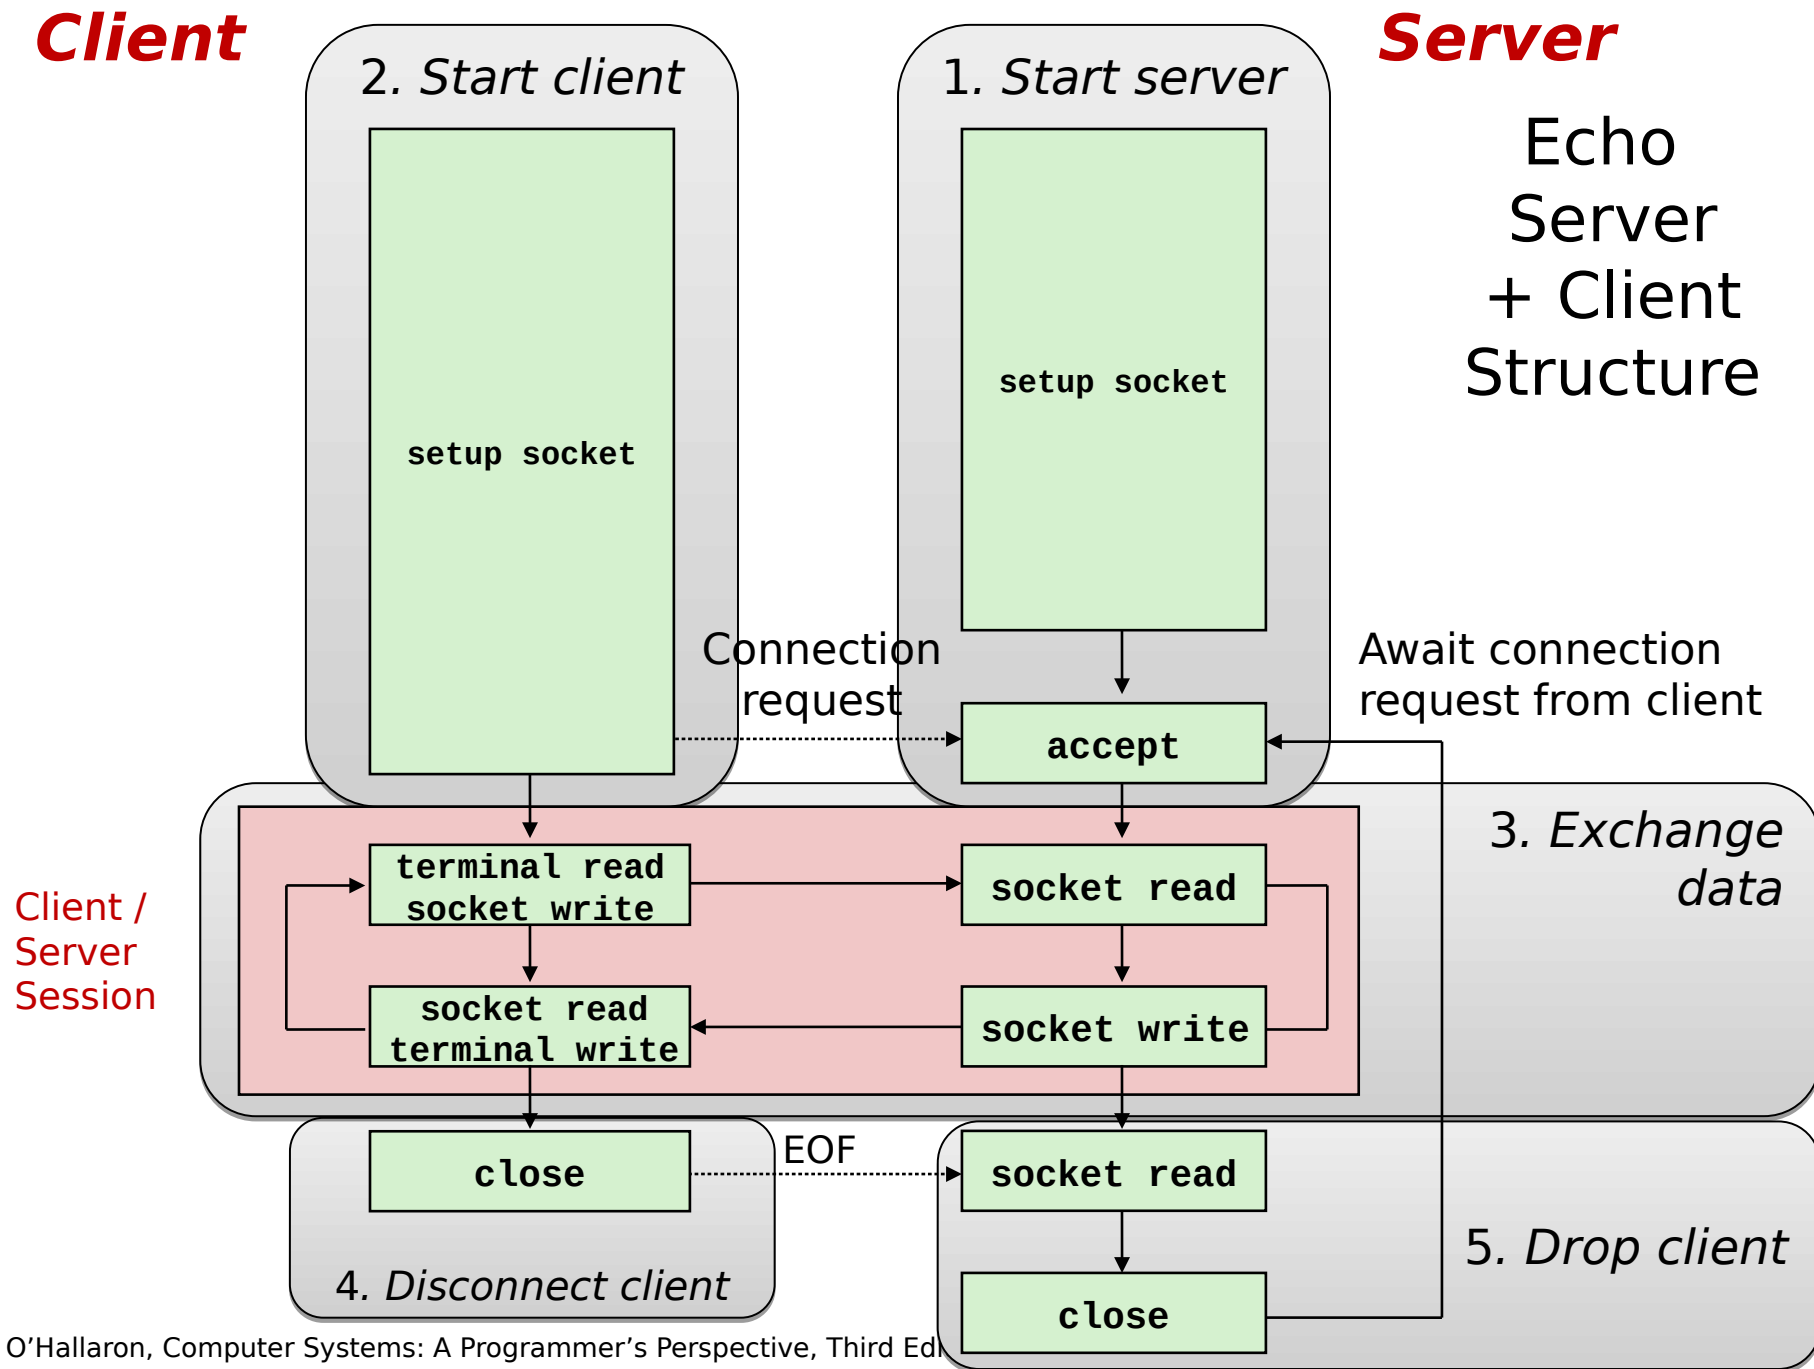
void freeaddrinfo(struct addrinfo *result); /* Free linked list */

const char *gai_strerror(int errcode);     /* Return error msg */
```

- Given **host** and **service**, **getaddrinfo** returns **result** that points to a linked list of **addrinfo** structs, each of which points to a corresponding socket address struct, and which contains arguments for the sockets interface functions.
- Helper functions:
 - **freeaddrinfo** frees the entire linked list.
 - **gai_strerror** converts error code to an error message.

Socket Programming Example

- Echo server and client
- **Server**
 - Accepts connection request
 - Repeats back lines as they are typed
- **Client**
 - Requests connection to server
 - Repeatedly:
 - Read line from terminal
 - Send to server
 - Read reply from server
 - Print line to terminal

Client**Server**

Echo Client: Main Routine

```
#include "compsys_helpers.h"

int main(int argc, char **argv)
{
    int clientfd;
    char *host, *port, buf[MAXLINE];
    compsys_helper_state_t state;

    host = argv[1];
    port = argv[2];

    clientfd = compsys_helper_open_clientfd(host, port);
    compsys_helper_readinitb(&state, clientfd);

    while (fgets(buf, MAXLINE, stdin) != NULL) {
        compsys_helper_writen(clientfd, buf, strlen(buf));
        compsys_helper_readlineb(&state, buf, MAXLINE);
        fputs(buf, stdout);
    }
    close(clientfd);
    exit(0);
}
```

echoclient.c

On Short Counts

- Short counts often occurs in these situations:
 - Encountering (end-of-file) EOF on reads
 - Reading text lines from a terminal
 - Reading and writing network sockets
- Short counts rarely occurs in these situations:
 - Reading from disk files (except for EOF)
 - ...but may happen for huge reads, depending on file system.
 - Writing to disk files
 - ...similarly.
- Best practice is to always allow for short counts.

The `compsys_helpers` package

- A set of wrappers that provide efficient and robust I/O in apps, such as network programs that are subject to short counts
- Provides two different kinds of functions
 - Unbuffered input and output of binary data
 - `compsys_helper_readn` and `compsys_helper_writen`
 - Buffered input of text lines and binary data
 - `compsys_helper_readlineb` and `compsys_helper_readnb`
 - Buffered functions are thread-safe and can be interleaved arbitrarily on the same descriptor
- Part of `compsys_helpers.c/compsys_helpers.h`
- For those that are resitting the course these are the same as the RIO functions. Feel free to continue using those but they are no longer supported

Unbuffered Compsys_helper

Input and Output

- Same interface as Unix `read` and `write`
- Especially useful for transferring data on network sockets

```
#include "compsys_helpers.h"
```

```
ssize_t compsys_helpers_readn(int fd, void *usrbuf, size_t n);  
ssize_t compsys_helpers_writen(int fd, void *usrbuf, size_t n);
```

**Return: num. bytes transferred if OK,
0 on EOF (compsys_helpers_readn only), -1 on error**

- `compsys_helpers_readn` returns short count only if it encounters EOF. Only use it when you know how many bytes to read
- `compsys_helpers_writen` never returns a short count
- Calls to `compsys_helpers_readn` and `compsys_helpers_writen` can be interleaved arbitrarily on the same descriptor

compsys_helper_readn

```
/*
 * compsys_helper_readn - Robustly read n bytes (unbuffered)
 */
ssize_t compsys_helper_readn(int fd, void *usrbuf, size_t n)
{
    size_t nleft = n;
    ssize_t nread;
    char *bufp = usrbuf;

    while (nleft > 0) {
        if ((nread = read(fd, bufp, nleft)) < 0) {
            if (errno == EINTR) /* Interrupted by sig handler return */
                nread = 0;      /* and call read() again */
            else
                return -1;      /* errno set by read() */
        }
        else if (nread == 0)
            break;              /* EOF */
        nleft -= nread;
        bufp += nread;
    }
    return (n - nleft);        /* Return >= 0 */
}
```

Buffered Input Helpers

- Efficiently read text lines and binary data from a file partially cached in an internal memory buffer

```
#include "compsys_helpers.h"

void compsys_helpers_readinitb(compsys_helper_state_t *rp, int fd);

ssize_t compsys_helpers_readlineb(
    compsys_helper_state_t *rp, void *usrbuf, size_t maxlen);
ssize_t compsys_helpers_readnb(
    compsys_helper_state_t *rp, void *usrbuf, size_t n);
```

Return: num. bytes read if OK, 0 on EOF, -1 on error

- **compsys_helpers_readlineb** reads a text line of up to **maxlen** bytes from file **fd** and stores the line in **usrbuf**. Especially useful for reading text lines from network sockets
 - Stopping conditions:
 - **maxlen** bytes read
 - EOF encountered
 - Newline (**'\n'**) encountered

Buffered Input Helpers

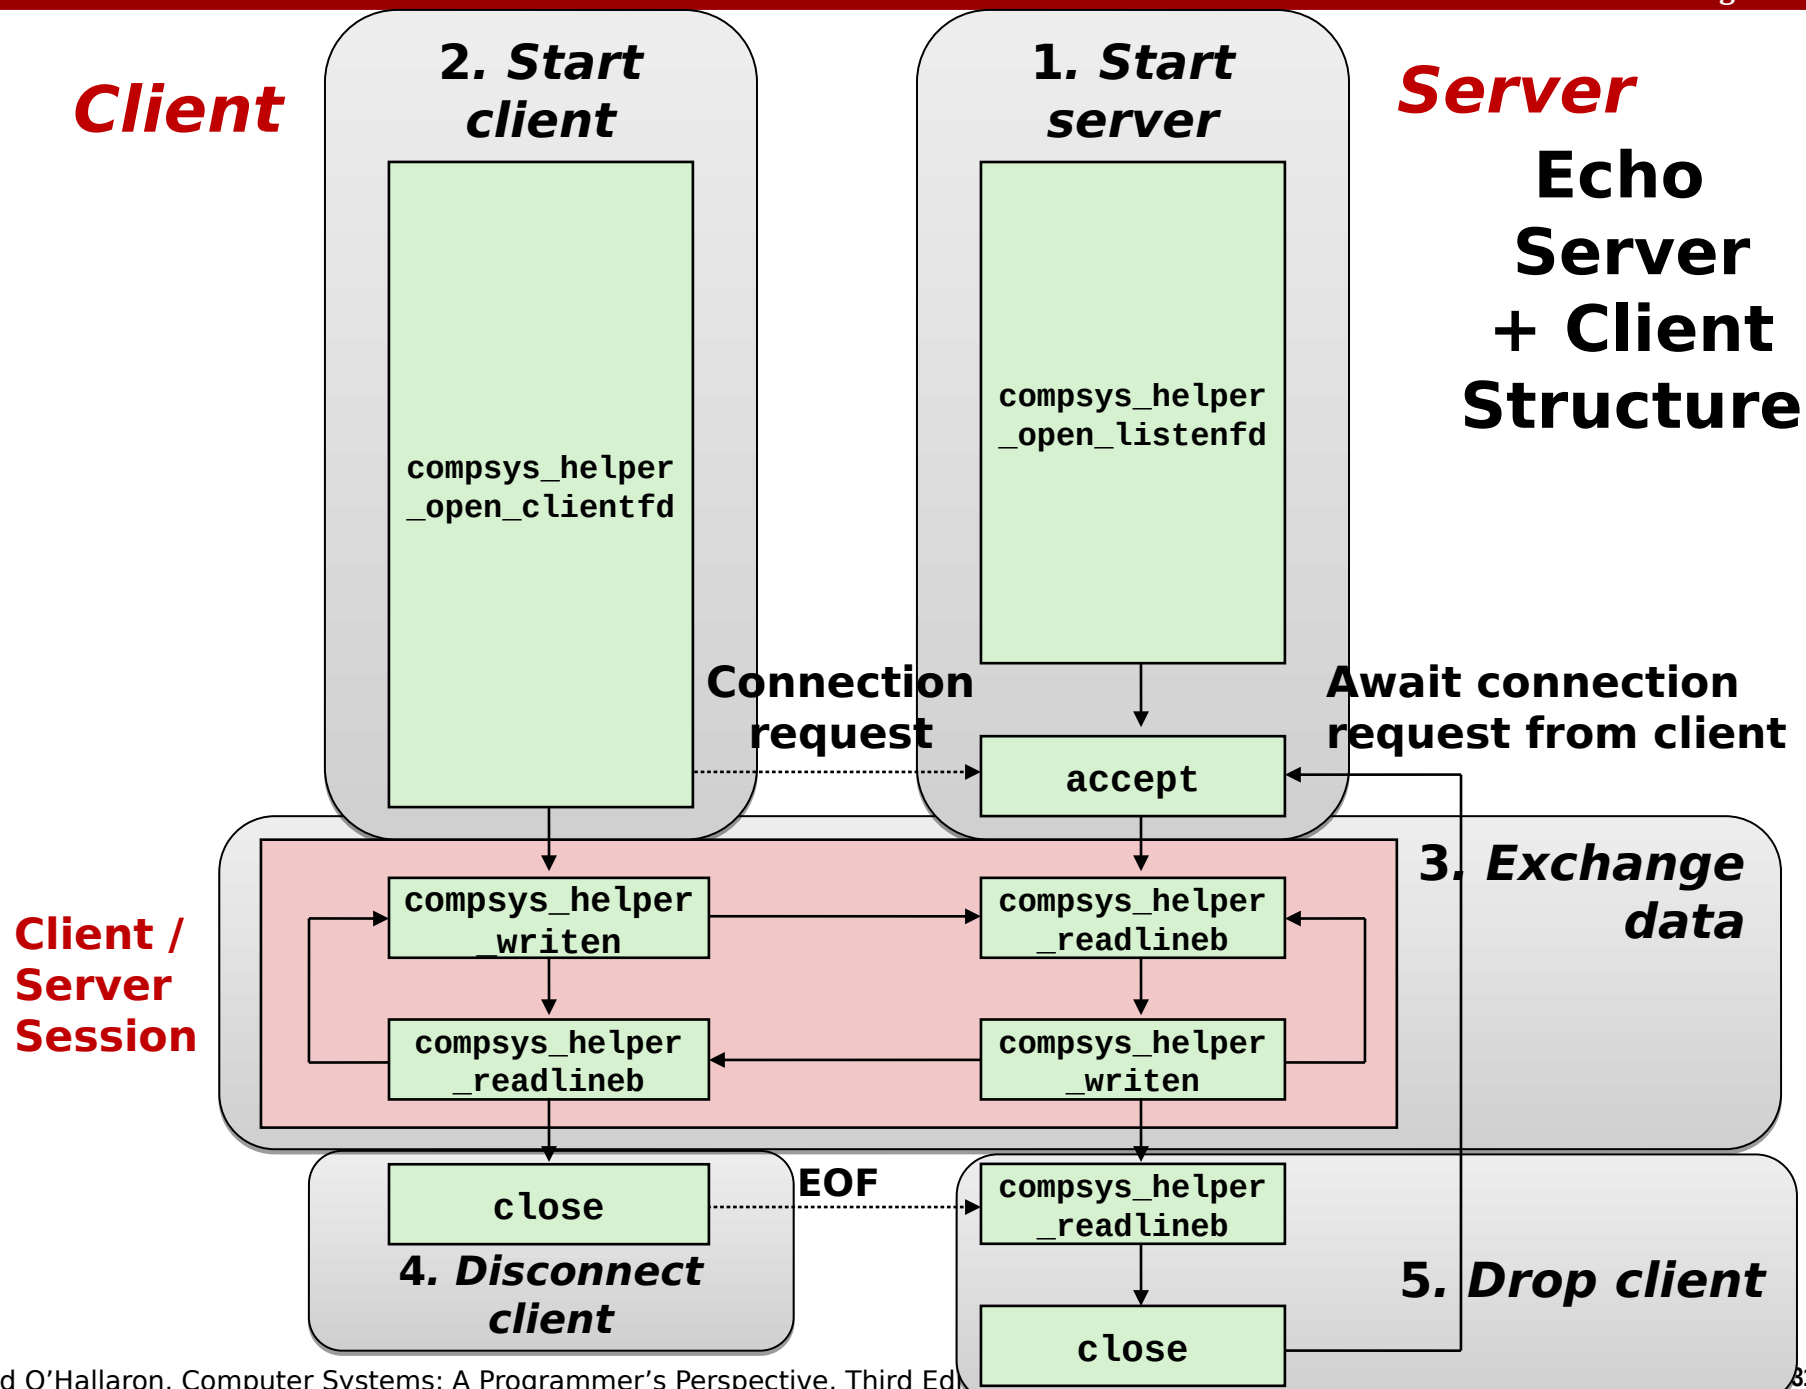
```
#include "compsys_helpers.h"

void compsys_helper_readinitb(compsys_helper_state_t *rp, int fd);

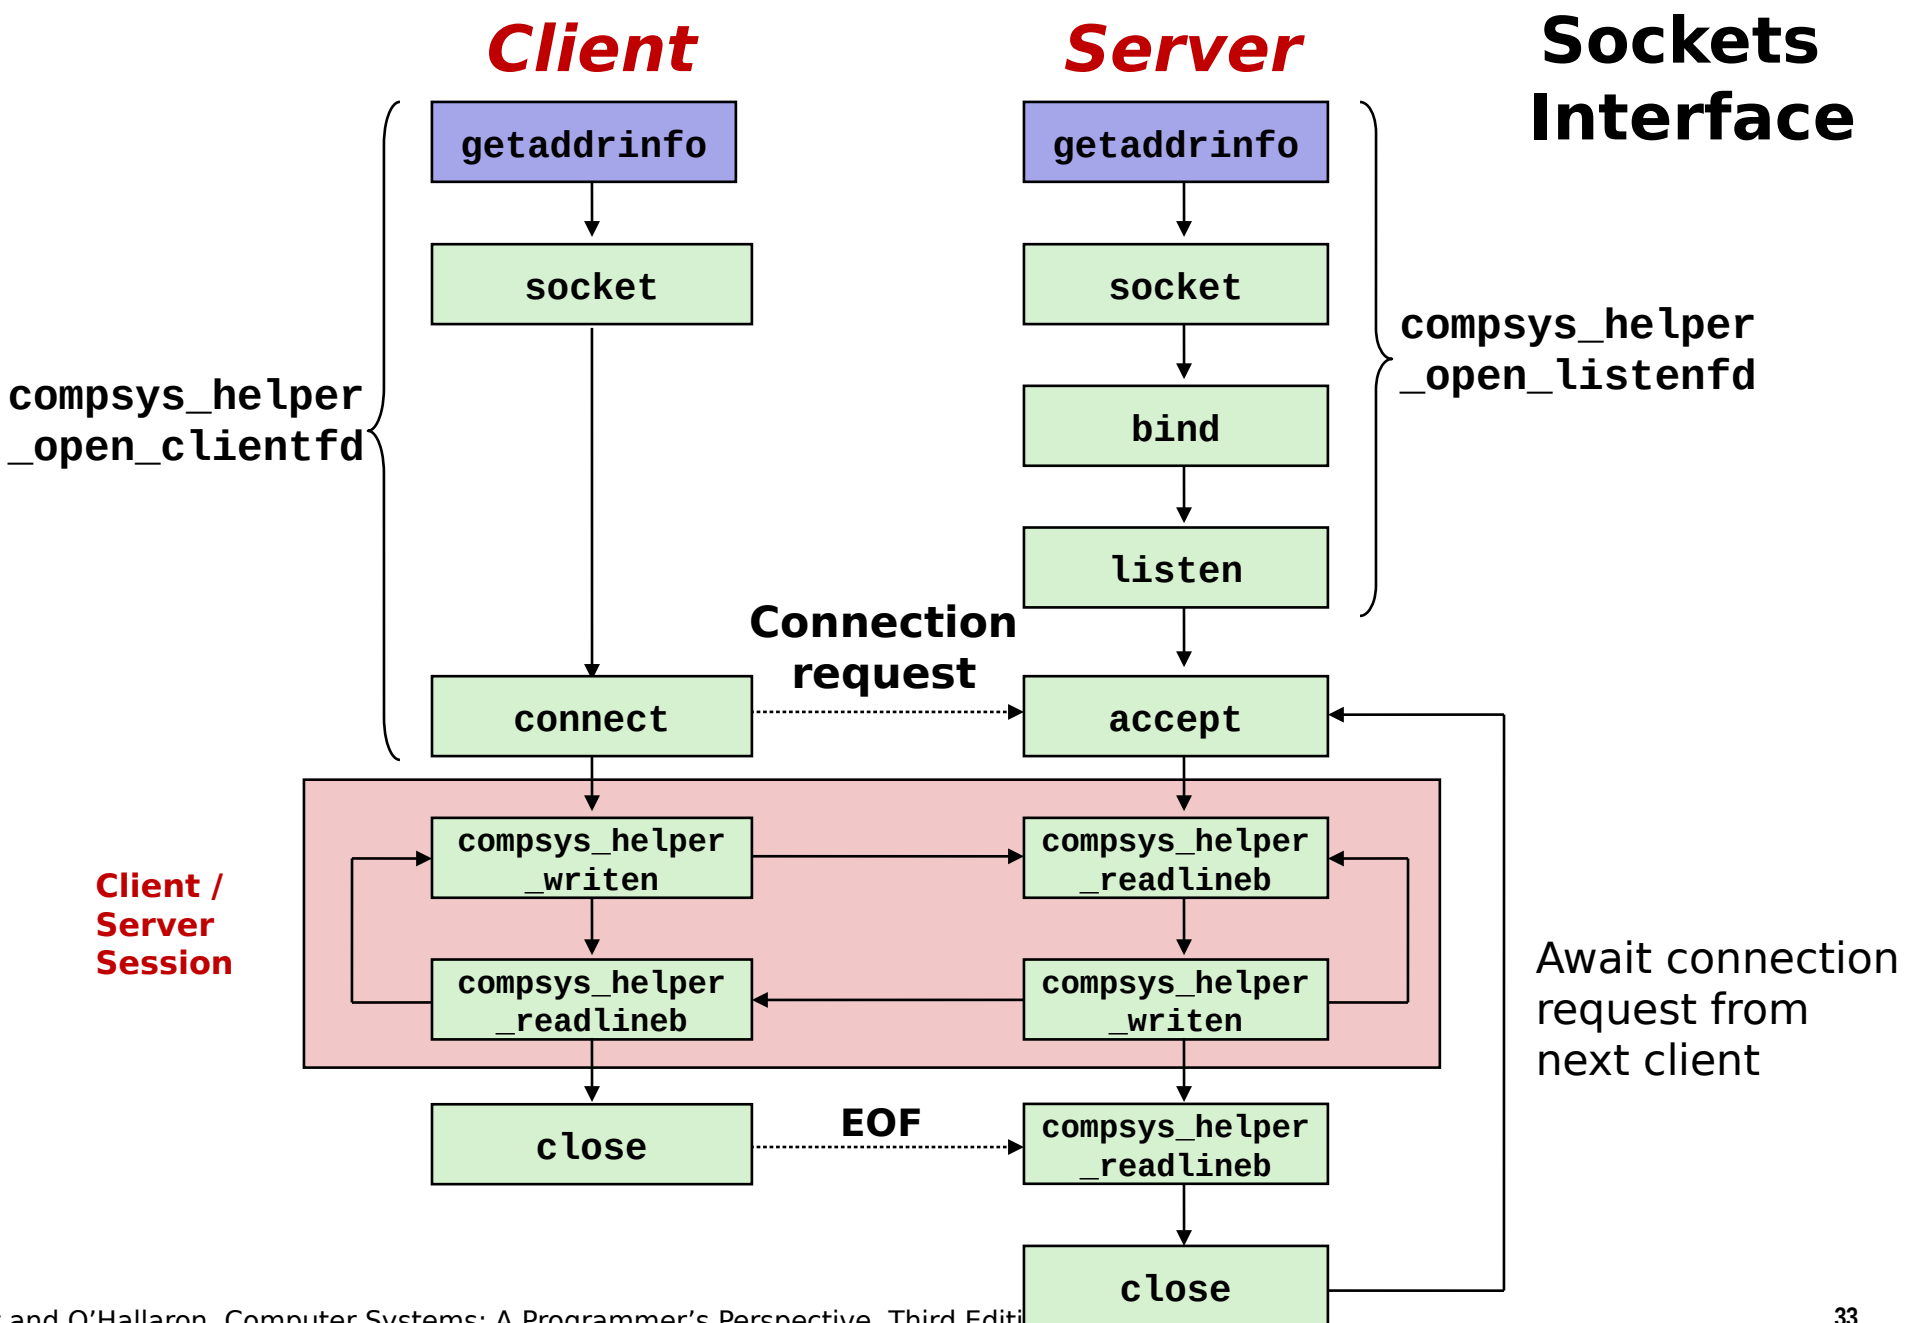
ssize_t compsys_helper_readlineb(
    compsys_helper_state_t *rp, void *usrbuf, size_t maxlen);
ssize_t compsys_helper_readnb(
    compsys_helper_state_t *rp, void *usrbuf, size_t n);
```

Return: num. bytes read if OK, 0 on EOF, -1 on error

- **compsys_helper_readnb** reads up to **n** bytes from file **fd**
- Stopping conditions
 - **maxlen** bytes read
 - EOF encountered
- Calls to **compsys_helper_readlineb** and **compsys_helper_readnb** can be interleaved arbitrarily on the same descriptor
 - Warning: Don't interleave with calls to **compsys_helper_readn**



See lecture code



Sockets Interface: socket

- Clients and servers use the socket function to create a *socket descriptor*:

```
int socket(int domain, int type, int protocol)
```

- Example:

```
int clientfd = socket(AF_INET, SOCK_STREAM, 0);
```

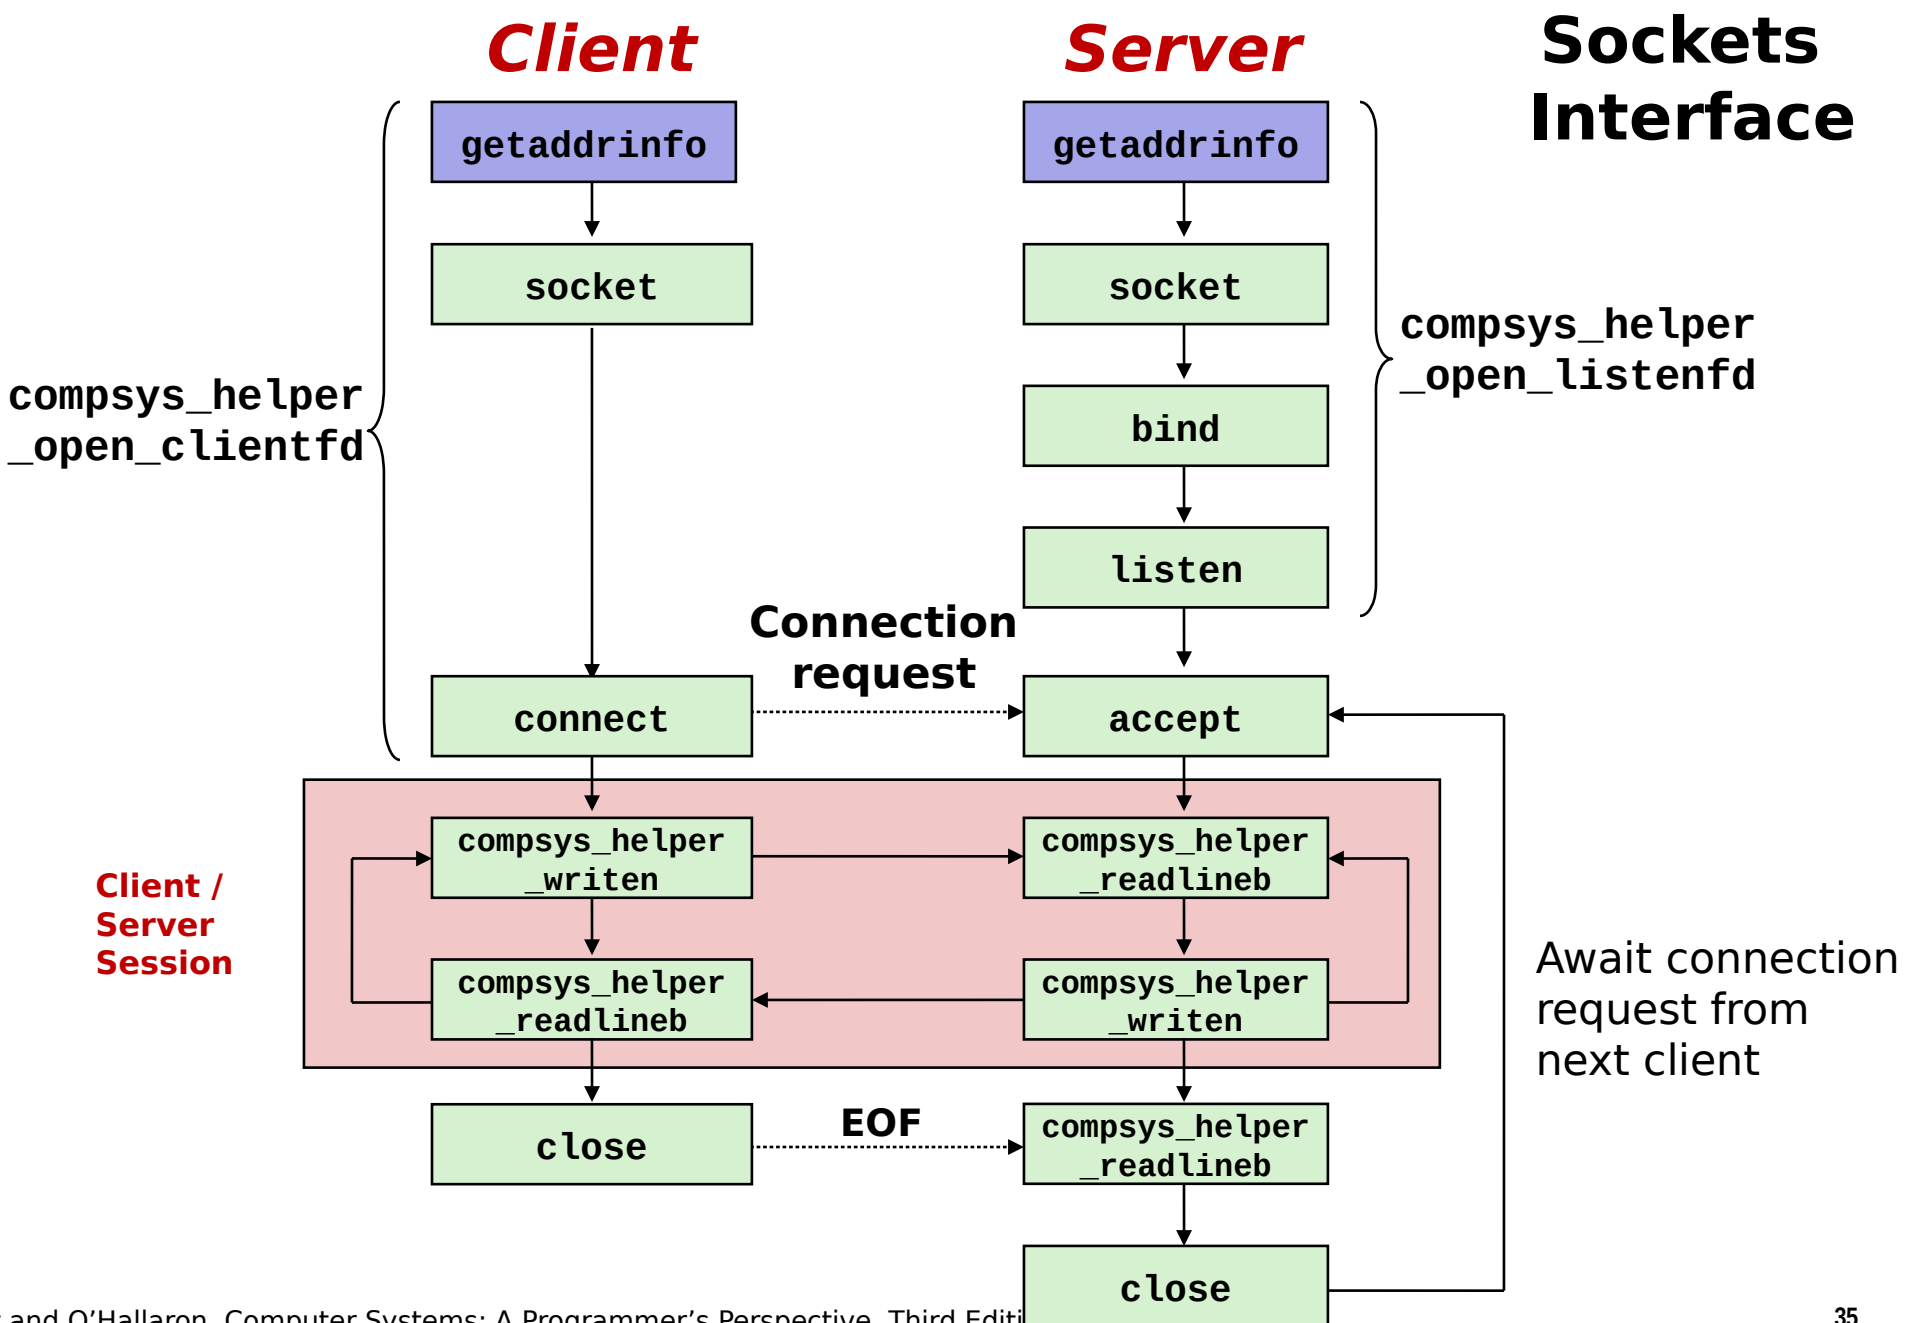


**Indicates that we
are using 32-bit
IPV4 addresses**



**Indicates that the
socket will be the
end point of a
connection**

Protocol specific! Best practice is to use `getaddrinfo` to generate the parameters automatically, so that code is protocol independent.



Sockets Interface: `bind`

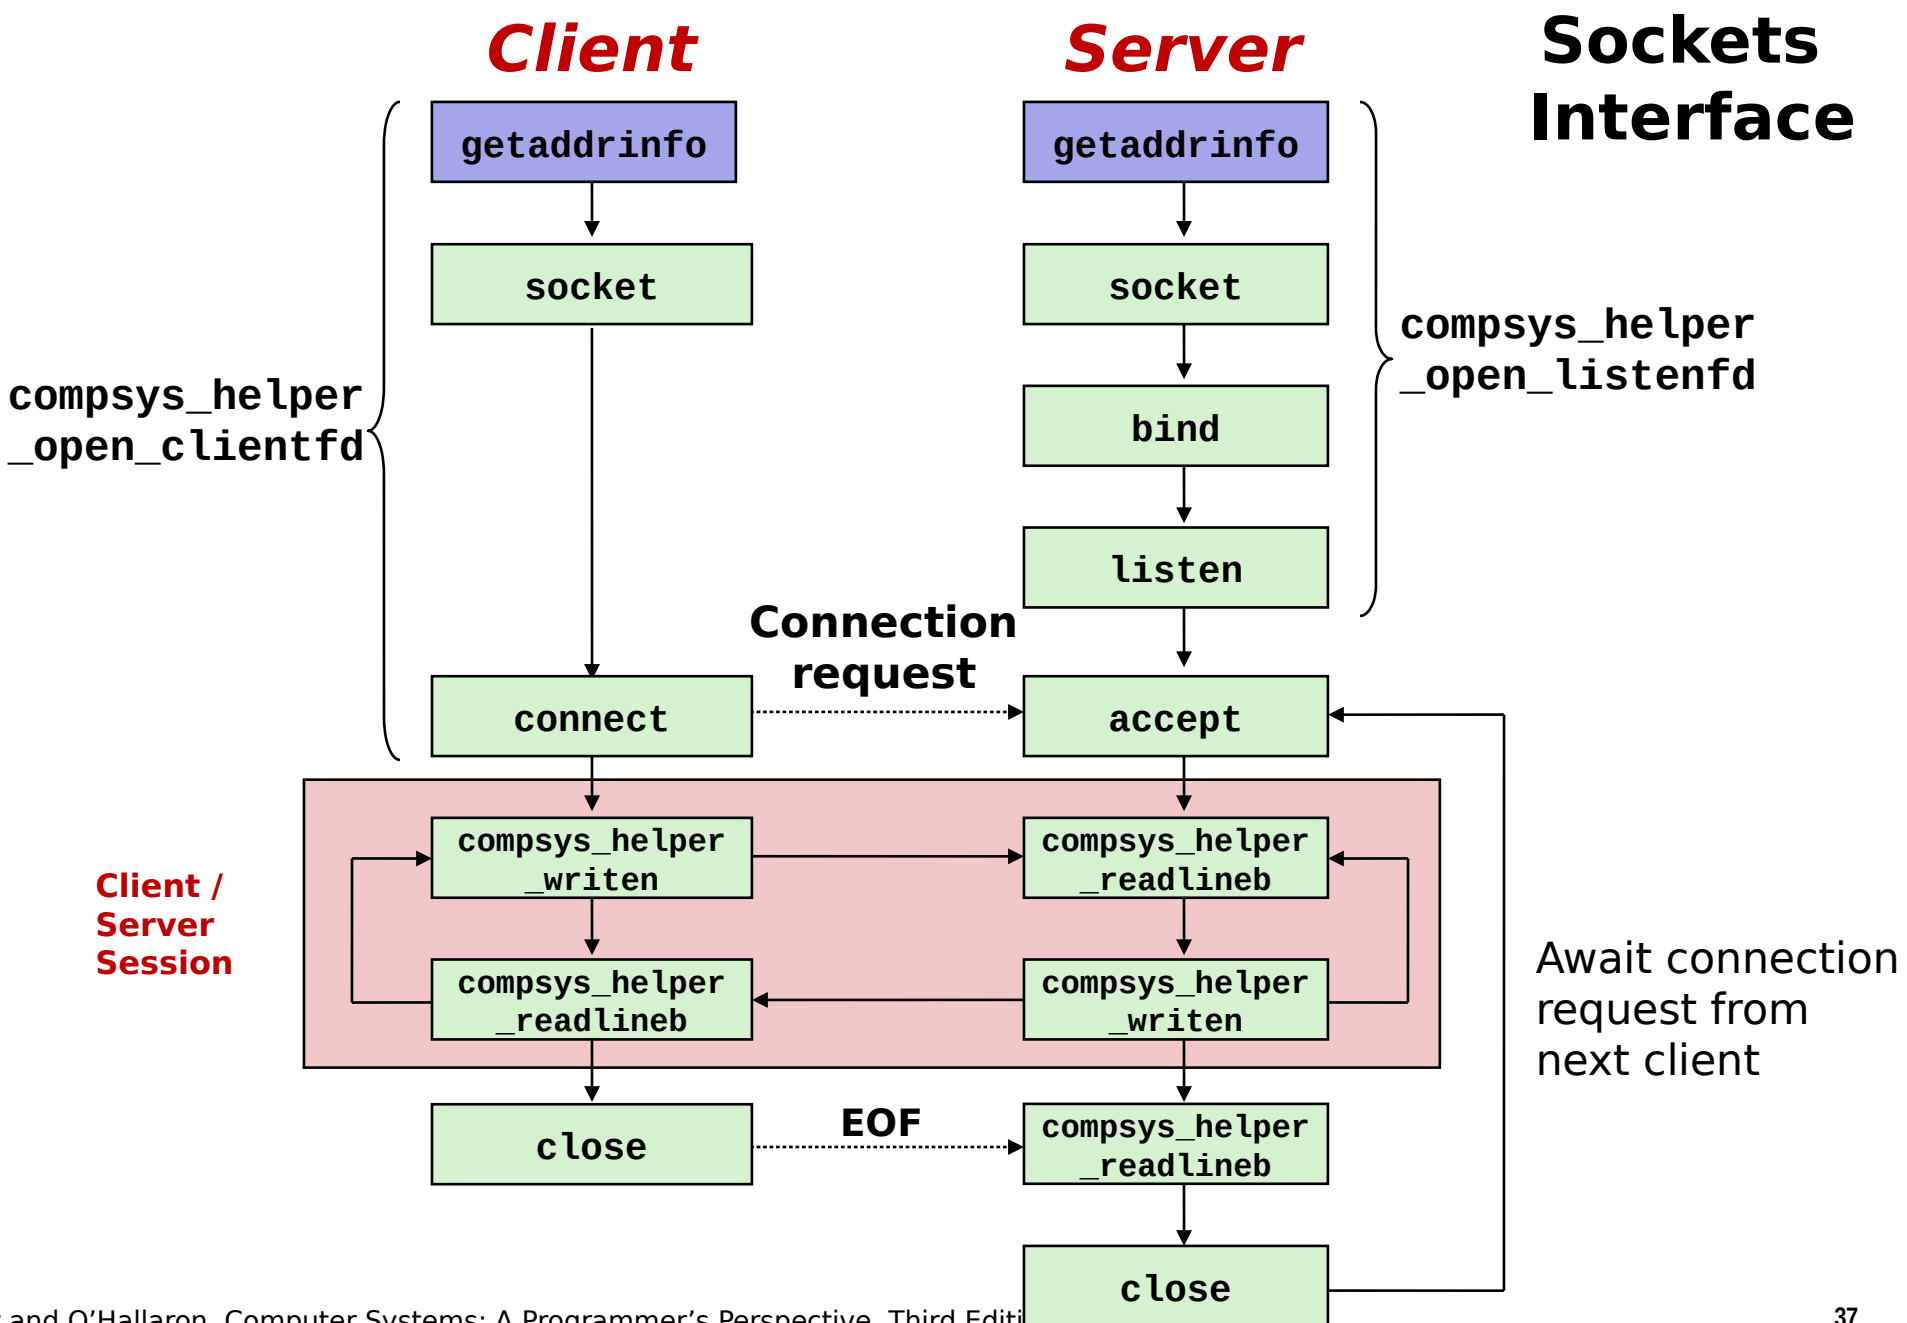
- A server uses `bind` to ask the kernel to associate the server's socket address with a socket descriptor:

```
int bind(int sockfd, SA *addr, socklen_t addrlen);
```

Recall: `typedef struct sockaddr SA;`

- Process can read bytes that arrive on the connection whose endpoint is `addr` by reading from descriptor `sockfd`
- Similarly, writes to `sockfd` are transferred along connection whose endpoint is `addr`

Best practice is to use `getaddrinfo` to supply the arguments `addr` and `addrlen`.

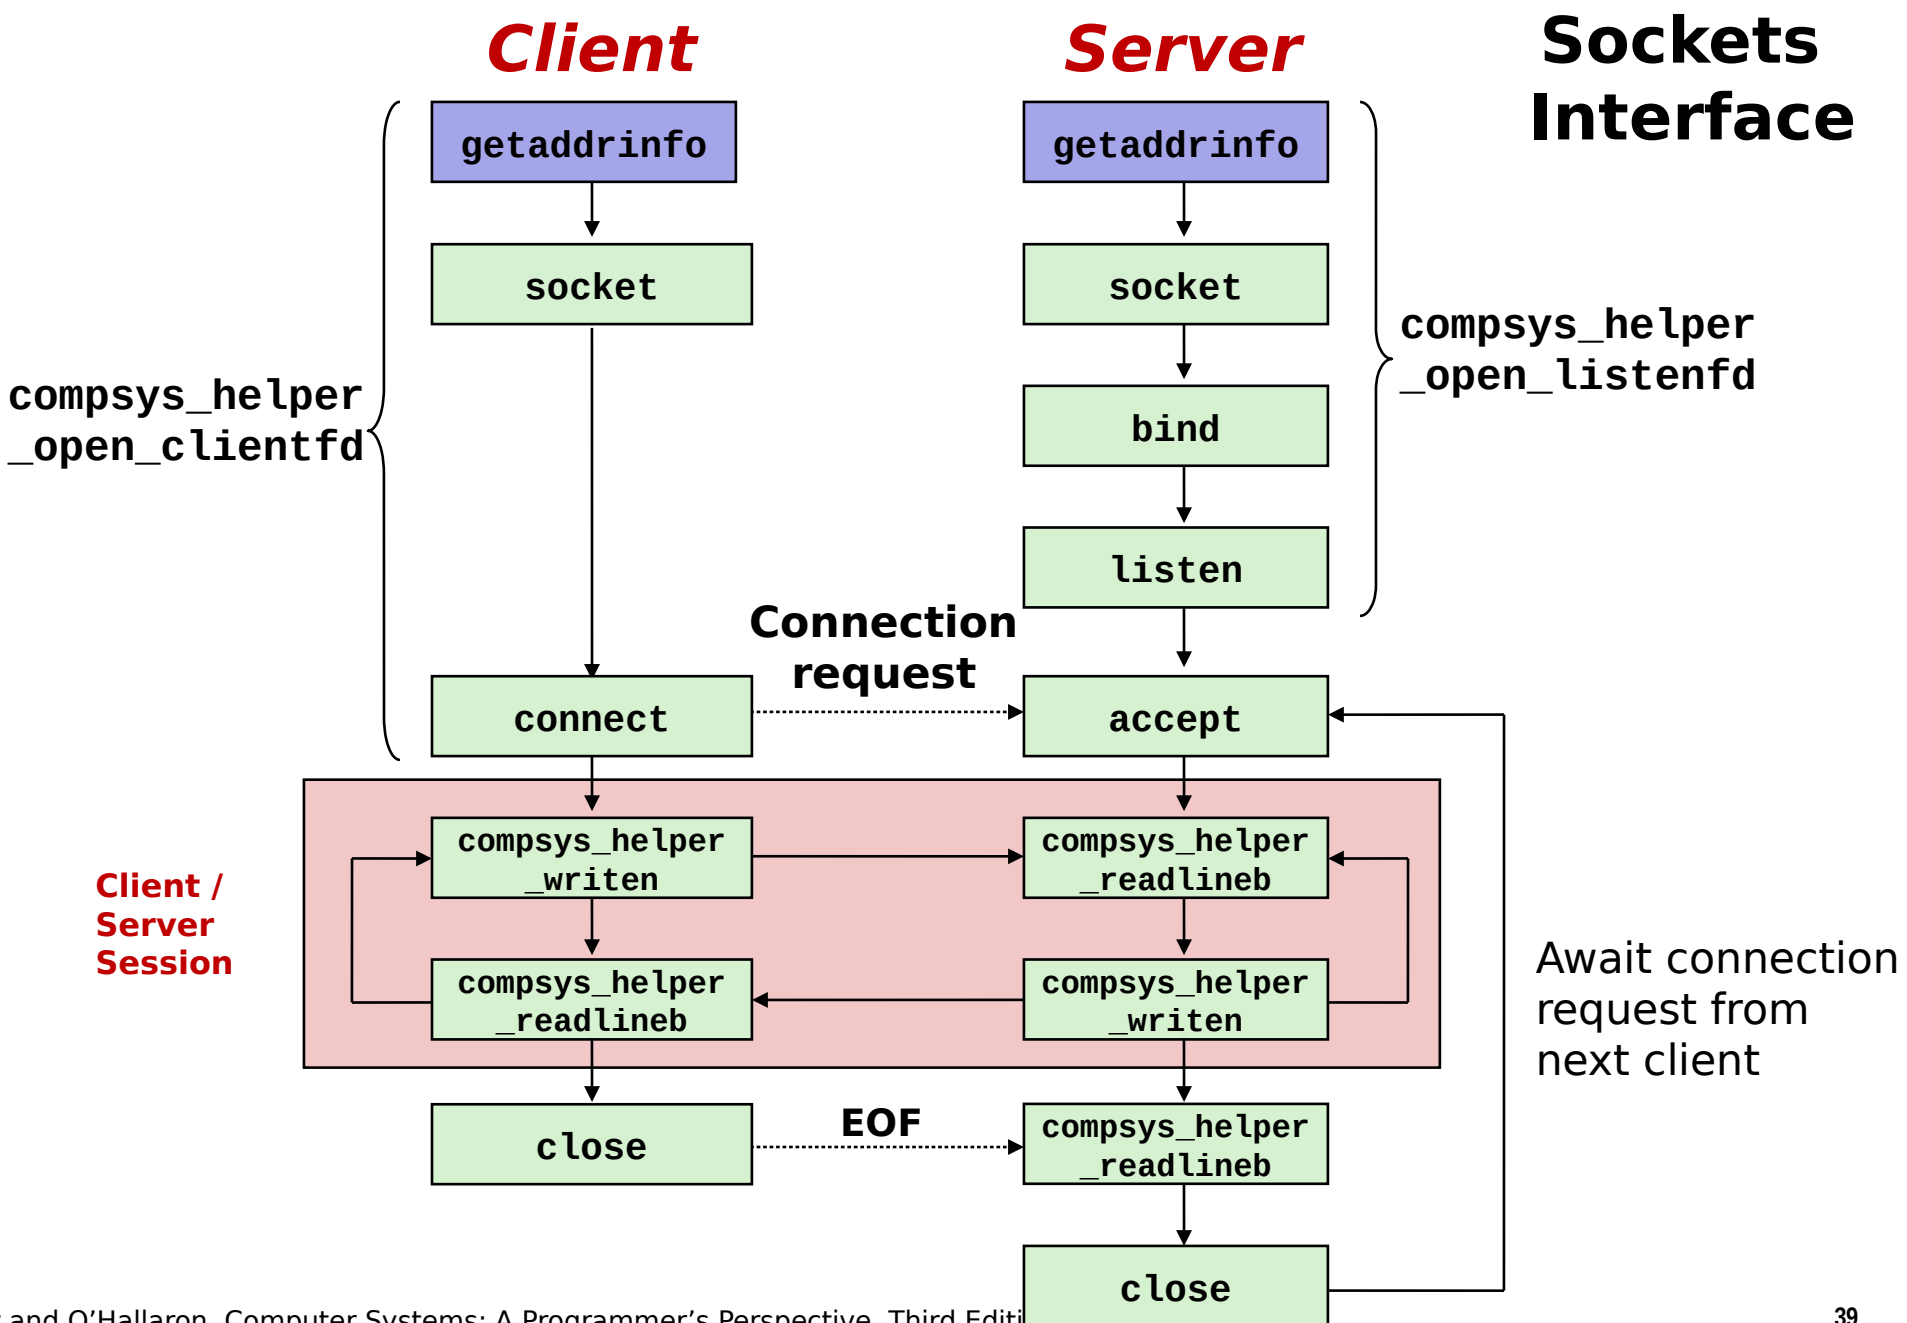


Sockets Interface: `listen`

- By default, kernel assumes that descriptor from `socket` function is an *active socket* that will be on the client end of a connection.
- A server calls the `listen` function to tell the kernel that a descriptor will be used by a server rather than a client:

```
int listen(int sockfd, int backlog);
```

- Converts **`sockfd`** from an active socket to a *listening socket* that can accept connection requests from clients.
- **`backlog`** is a hint about the number of outstanding connection requests that the kernel should queue up before starting to refuse requests.

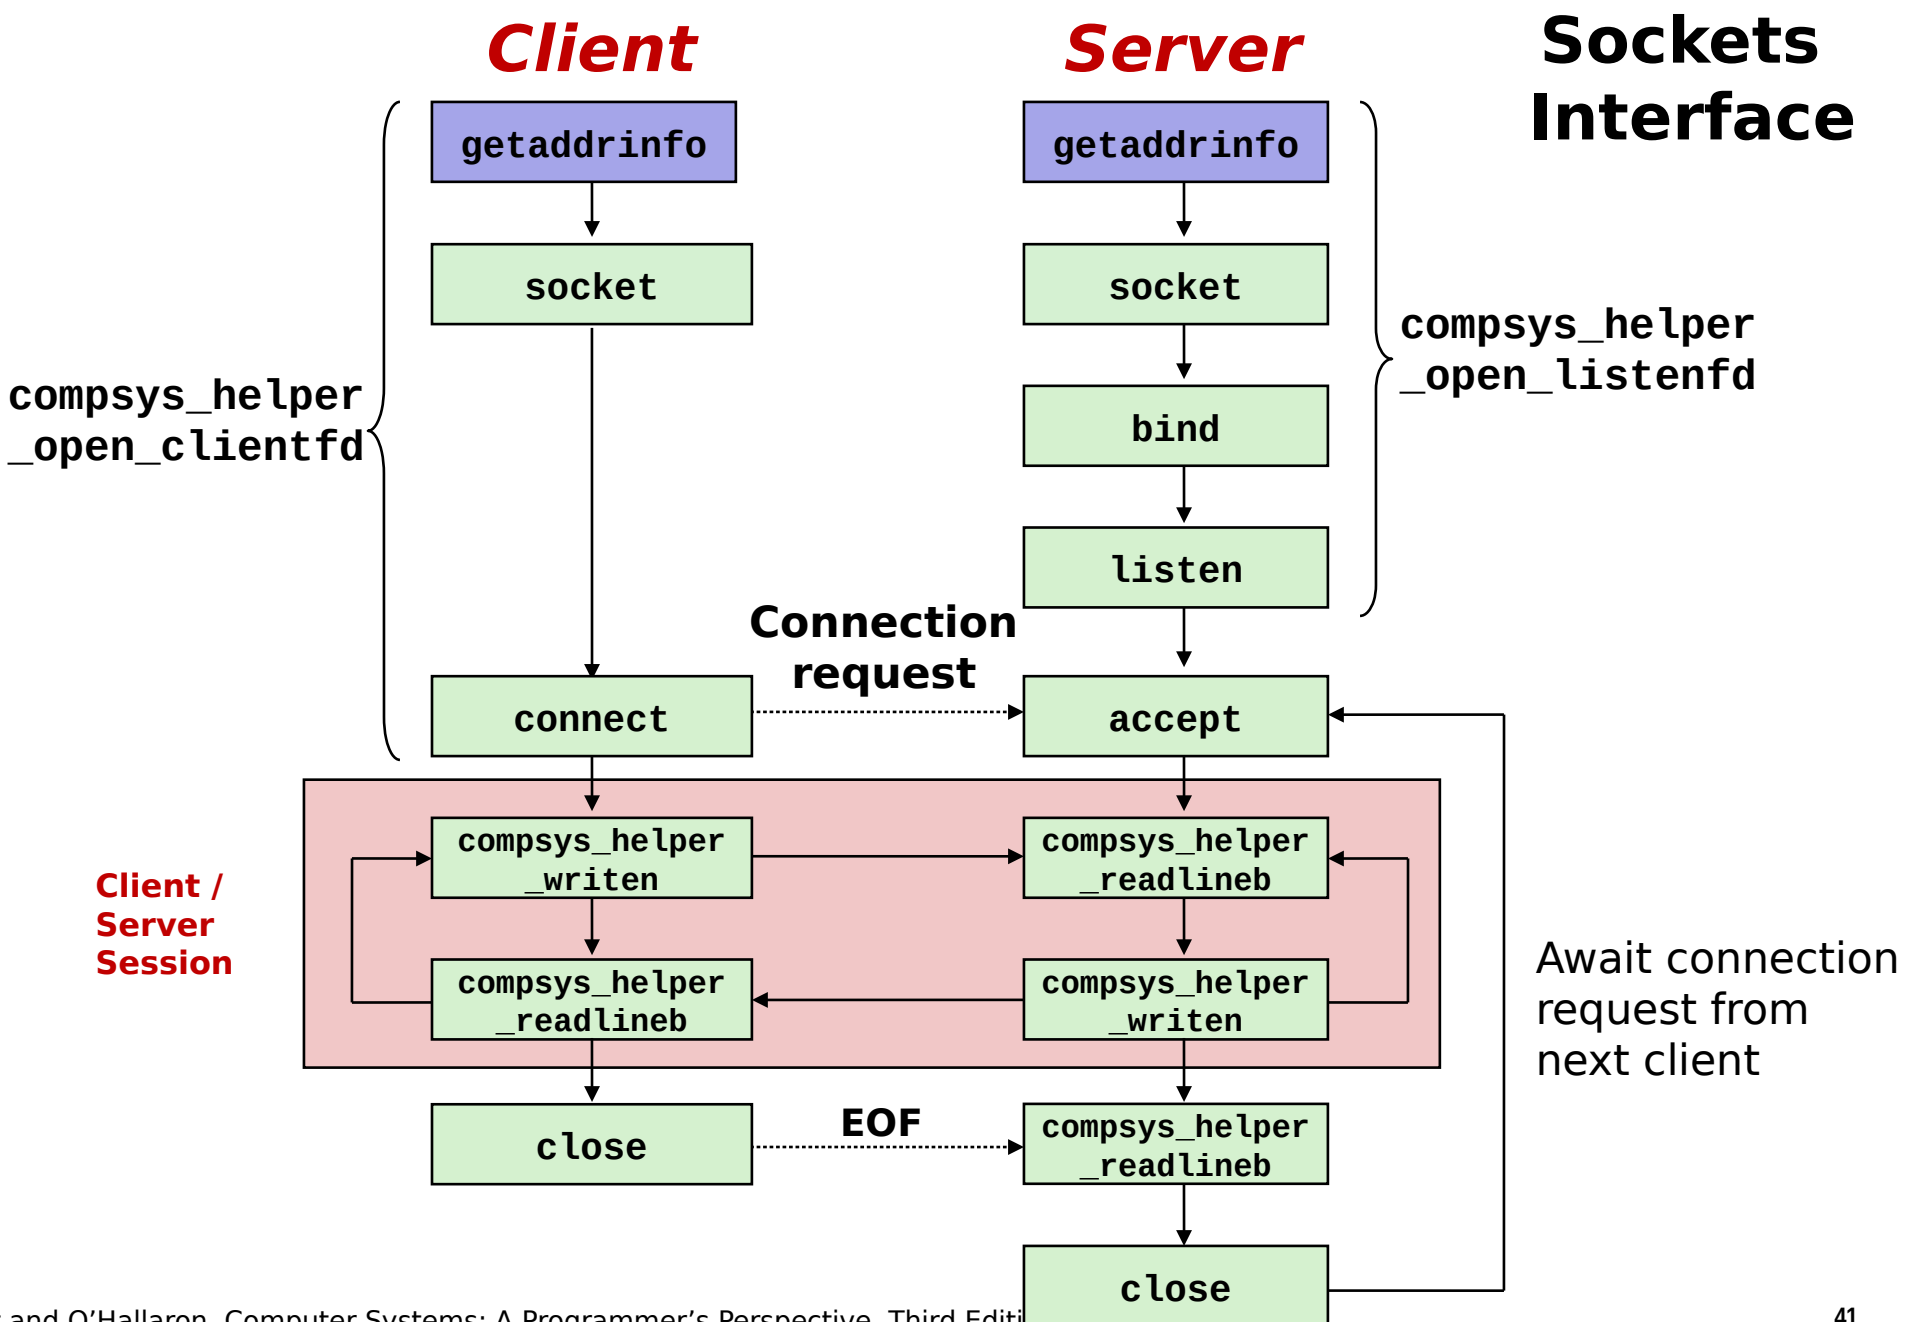


Sockets Interface: accept

- Servers wait for connection requests from clients by calling **accept**:

```
int accept(int listenfd, SA *addr, int *addrlen);
```

- Waits for connection request to arrive on the connection bound to **listenfd**, then fills in client's socket address in **addr** and size of the socket address in **addrlen**.
- Returns a *connected descriptor* that can be used to communicate with the client via Unix I/O routines.



Sockets Interface: connect

- A client establishes a connection with a server by calling connect:

```
int connect(int clientfd, SA *addr, socklen_t addrlen);
```

- Attempts to establish a connection with server at socket address **addr**
 - If successful, then `clientfd` is now ready for reading and writing.
 - Resulting connection is characterized by socket pair `(x:y, addr.sin_addr:addr.sin_port)`
 - `x` is client address
 - `y` is ephemeral port that uniquely identifies client process on client host

Best practice is to use **getaddrinfo** to supply the arguments **addr** and **addrlen**.

accept Illustrated



1. Server blocks in ***accept***, waiting for connection request on listening descriptor ***listenfd***



2. Client makes connection request by calling and blocking in ***connect***



3. Server returns ***connfd*** from ***accept***. Client returns from ***connect***. Connection is now established between ***clientfd*** and ***connfd***

Connected vs. Listening Descriptors

■ Listening descriptor

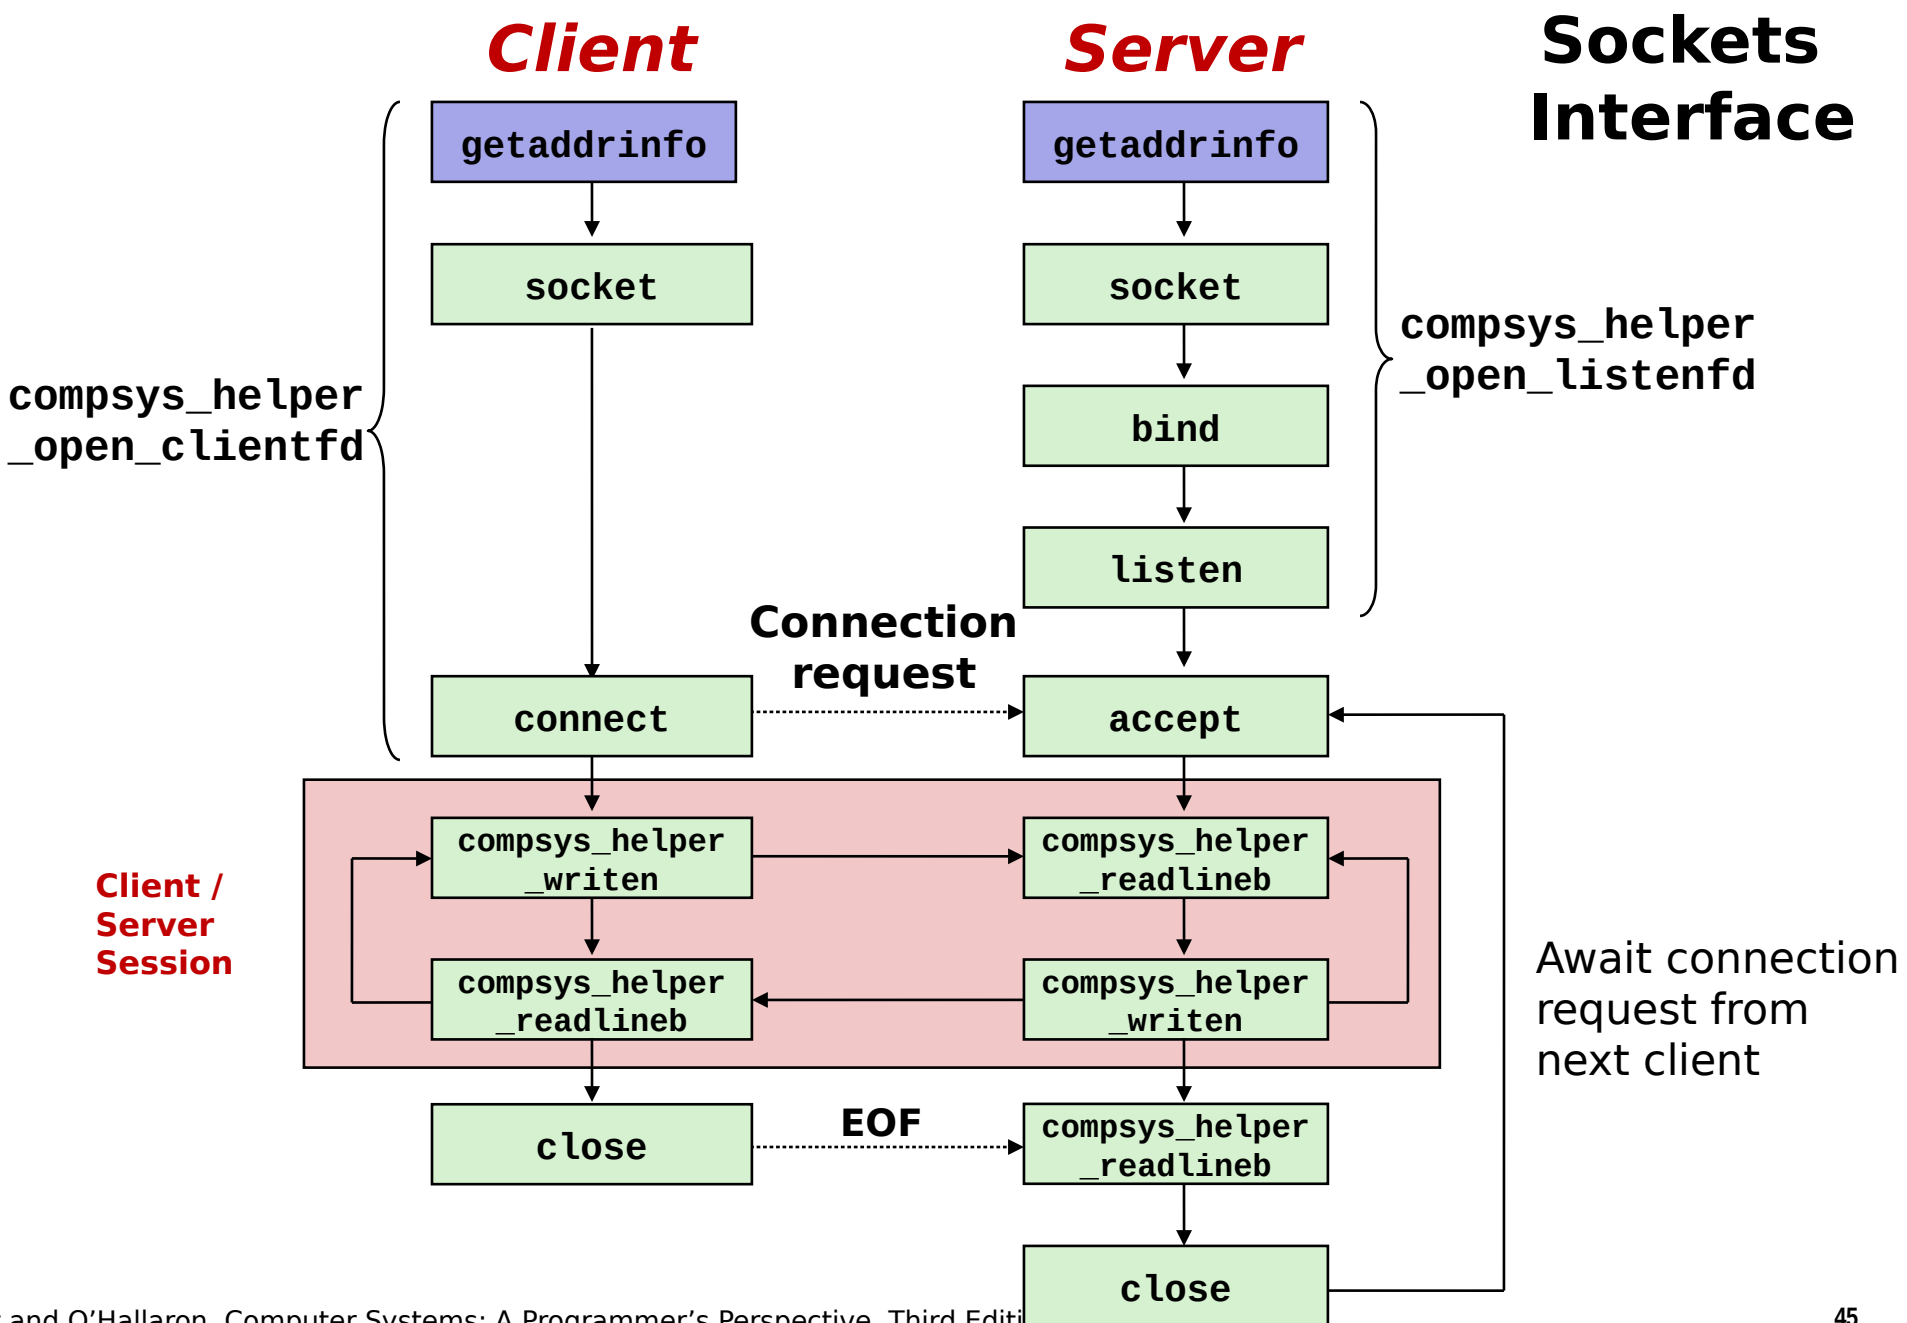
- End point for client connection requests
- Created once and exists for lifetime of the server

■ Connected descriptor

- End point of the connection between client and server
- A new descriptor is created each time the server accepts a connection request from a client
- Exists only as long as it takes to service client

■ Why the distinction?

- Allows for concurrent servers that can communicate over many client connections simultaneously
 - E.g., Each time we receive a new request, we fork a child to handle the request



Python for Networks

- **Higher level than C, so should be easier to follow and understand**
- **More abstractions, so quicker to get a working networked application, but runs slower**
- **Typically, you are more likely to use it yourselves so its worth introduction. Assignments will still be in C (sorry not sorry)**



```
def function(num):  
    for i in [1, 2, 3, 4]:  
        print(num + i)  
  
    return num * 2  
  
print(function(10))
```

Sockets: socket

- Clients and servers use the socket function to create a socket descriptor:

```
int socket(int domain, int type, int protocol)
```

- Example:

C:

```
#include <sys/socket.h>
```

```
int socket_fd = socket(AF_INET, SOCK_STREAM, 0);
```

↑
**Indicates that we
are using 32-bit
IPV4 addresses**

↑
**Indicates that the
socket will be the end
point of a connection**

Python:

```
from socket import *  
with socket(AF_INET, SOCK_STREAM) as sock:  
    ...
```



Sockets: listen

- By default, kernel assumes that descriptor from socket function is an active socket that will be on the client end of a connection.
- A server calls the listen function to tell the kernel that a descriptor will be used by a server rather than a client:

```
int listen(int sockfd, int backlog);
```

- Converts sockfd from an active socket to a listening socket that can accept connection requests from clients.

C: `listen(socket_fd, 10);`

Python: `sock.listen(10)`



Sockets: accept

- Servers wait for connection requests from clients by calling accept:

```
int accept(int listenfd, SA *addr, int *addrlen);
```

- Waits for connection request to arrive on the connection bound to listenfd, then fills in client's socket address in addr and size of the socket address in addrlen.
- Returns a connected descriptor that can be used to communicate with the client via Unix I/O routines.

C:

```
socklen_t clientlen;  
struct sockaddr_storage clientaddr;  
conn_fd = accept(socket_fd, (SA *) &clientaddr, &clientlen);
```

Python:

```
Conn, conn_addr = sock.accept()
```



Sockets: connect

- A client establishes a connection with a server by calling connect:

```
int connect(int clientfd, SA *addr, socklen_t addrlen);
```

- Attempts to establish a connection with server at socket address addr
 - If successful, then clientfd is now ready for reading and writing.
 - Resulting connection is characterized by socket pair (x:y, addr.sin_addr:addr.sin_port)
 - x is client address
 - y is ephemeral port that uniquely identifies client process on client host

C:

```
struct sockaddr s_addr;  
connect(socket_fd, (struct sockaddr *)&s_addr, sizeof(s_addr));
```

Python:

```
client_sock.connect("130.226.237.173", 56)
```



Final building blocks

- Reading from Python socket:

```
socket.recv(buffsize)
```

- Writing to Python socket:

```
socket.send(bytes)
```

```
socket.sendall(bytes)
```

- Both send bytes, but send may only send some and it is your responsibility to check. Sendall manages sending until everything is sent or an error was encountered



Python Example

Client:

```
import socket

with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as client_socket:
    client_socket.connect(("127.0.0.1", 5678))
    request = bytearray("This is a message".encode())
    client_socket.sendall(request)
    response = client_socket.recv(1024)
    print(response)
```

Server:

```
import socket

with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as server_socket:
    server_socket.bind(("127.0.0.1", 5678))
    server_socket.listen()
    while True:
        connection, connection_address = server_socket.accept()
        with connection:
            message = connection.recv(1024)
            connection.sendall(response)
```

Bytes in Python

- In networking we need to be deliberate in what bytes we send, but Python does not like operating at this level
- Bytearrays must be manually packed and extended:

```
import struct

payload = bytearray()
payload.extend("Some long string.")
payload.extend(4798.5)
payload.extend(struct.pack('!I', 4294967295))
payload.extend(struct.pack('!I', 0))
```

- Key difference:
 - Extend will simply add its input to the end of the array, usefull for message bodies
 - struct.pack takes a formatting variable defining exactly how much space a variable should take up, and the endianness of the bytes
 - formatting: <https://docs.python.org/3/library/struct.html>



Summary

- **Sockets used to communicate across processes over a network (even same network card)**
 - TCP sockets – Listening vs connecting sockets
 - Quirks in structs representing network addresses.
 - Use getaddrinfo() or fill up the struct yourself.
 - Usage of compsys_helpers library for buffered I/O.

Testing Servers Using telnet

- **The telnet program is invaluable for testing servers that transmit ASCII strings over Internet connections**
 - Our simple echo server
 - Web servers
 - Mail servers

- **Usage:**
 - `linux> telnet <host> <portnumber>`
 - Creates a connection with a server running on *<host>* and listening on port *<portnumber>*